

## Assignment 2

### Big Data Technologies – CS554

Prof. Navendu Garg

ID : A20522183

Pem key created, on the EC2 instance



Created EMR – cluster on the console.

Clone
Terminate
AWS CLI export

Cluster: **asn3-cluster** Waiting Cluster ready to run steps.

Summary
Application user interfaces
Monitoring
Hardware
Configurations
Events
Steps
Bootstrap actions

#### Summary

ID: j-13i8LJ4U4MTB4

Creation date: 2023-09-21 00:38 (UTC-5)

Elapsed time: 7 minutes

After last step completes: Cluster waits

Termination protection: Off [Change](#)

Tags: -- [View All](#) / [Edit](#)

Master public DNS: ec2-34-228-195-135.compute-1.amazonaws.com [🔗](#)

[Connect to the Master Node Using SSH](#)

#### Application user interfaces

Persistent user interfaces [🔗](#): [YARN timeline server](#), [Tez UI](#)

On-cluster user interfaces [🔗](#): Not Enabled [Enable an SSH Connection](#)

#### Configuration details

Release label: emr-5.36.1

Hadoop distribution: Amazon 2.10.1

Applications: Hive 2.3.9, Hue 4.10.0, Mahout 0.13.0, Pig 0.17.0, Tez 0.9.2

Log URI: s3://aws-logs-965044767171-us-east-1/elasticmapreduce/ [📁](#)

EMRFS consistent view: Disabled

Custom AMI ID: --

Amazon Linux Release: 2.0.20230822.0 [Learn more](#) [🔗](#)

#### Network and hardware

Availability zone: us-east-1b

Subnet ID: [subnet-0a2b5c65c3df58ce0](#) [🔗](#)

Master: **Bootstrapping** 1 m5.xlarge

Core: **Provisioning** 1 m5.xlarge

Task: --

Cluster scaling: Not enabled

Auto-termination: Terminate if idle for 1 day



Sudo install /usr/bin/pip3.7 install mrjob[aws].

The above command installs mrjob libraries and services for the hadoop map reduce for python files.

```
hadoop@ip-172-31-21-101:~$ sudo /usr/bin/pip3.7 install mrjob[aws]
WARNING: Running pip install with root privileges is generally not a good idea.
Try `pip3.7 install --user` instead.
Collecting mrjob[aws]
  Downloading mrjob-0.7.4-py2.py3-none-any.whl (439 kB)
    | 439 kB 25.7 MB/s
Requirement already satisfied: PyYAML>=3.10 in /usr/local/lib64/python3.7/site-packages (from mrjob[aws]) (5.4.1)
Collecting boto3>=1.13.26; extra == "aws"
  Downloading boto3-1.31.52-py3-none-any.whl (11.2 MB)
    | 11.2 MB 68.5 MB/s
Collecting boto3>=1.10.0; extra == "aws"
  Downloading boto3-1.28.52-py3-none-any.whl (135 kB)
    | 135 kB 78.6 MB/s
Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in /usr/local/lib/python3.7/site-packages (from boto3>=1.13.26; extra == "aws"->mrjob[aws]) (1.0.0)
Collecting urllib3<1.27,>=1.25.4
  Downloading urllib3-1.26.16-py2.py3-none-any.whl (143 kB)
    | 143 kB 75.2 MB/s
Collecting python-dateutil<3.0.0,>=2.1
  Downloading python_dateutil-2.8.2-py2.py3-none-any.whl (247 kB)
    | 247 kB 79.9 MB/s
Collecting s3transfer<0.7.0,>=0.6.0
```

Successfully connected to pip python3.7

```
hadoop@ip-172-31-21-101:~$ sudo /usr/bin/pip3.7 install mrjob[aws]
Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in /usr/local/lib/python3.7/site-packages (from boto3>=1.13.26; extra == "aws"->mrjob[aws]) (1.0.0)
Collecting urllib3<1.27,>=1.25.4
  Downloading urllib3-1.26.16-py2.py3-none-any.whl (143 kB)
    | 143 kB 75.2 MB/s
Collecting python-dateutil<3.0.0,>=2.1
  Downloading python_dateutil-2.8.2-py2.py3-none-any.whl (247 kB)
    | 247 kB 79.9 MB/s
Collecting s3transfer<0.7.0,>=0.6.0
  Downloading s3transfer-0.6.2-py3-none-any.whl (79 kB)
    | 79 kB 19.4 MB/s
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil<3.0.0,>=2.1->boto3>=1.13.26; extra == "aws"->mrjob[aws]) (1.13.0)
Installing collected packages: urllib3, python-dateutil, boto3, s3transfer, mrjob
  WARNING: The scripts mrjob, mrjob-3 and mrjob-3.7 are installed in '/usr/local/bin' which is not on PATH.
  Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
Successfully installed boto3-1.28.52 botocore-1.31.52 mrjob-0.7.4 python-dateutil-2.8.2 s3transfer-0.6.2 urllib3-1.26.16
hadoop@ip-172-31-21-101:~$
```

```
Successfully installed boto3-1.28.52 botocore-1.31.52 mrjob-0.7.4 python-dateutil-2.8.2 s3transfer-0.6.2 urllib3-1.26.16
[hadoop@ip-172-31-21-101 ~]$
```

#### Step 4

Word count py , is securely copied using the command `scp -i /path/to/.pem /path/to/wordcount.py`

```
saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ scp -i C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem C:/Users/saira/OneDrive/Desktop/WordCount.py hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoop
WordCount.py                                100% 402      8.8KB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$
```

`scp -i /path/to/.pem /path/to/w.data`

```
MINGW64:/c:/Users/saira/OneDrive/Desktop
[-i identity_file] [-J destination] [-l limit]
[-o ssh_option] [-P port] [-S program] source ... target

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ scp -i C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem C:/Users/saira/OneDrive/Desktop/WordCount.py hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoop
scp: dest open "home/hadoop": No such file or directory
scp: failed to upload file C:/Users/saira/OneDrive/Desktop/WordCount.py to home/hadoop

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ scp -i C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem C:/Users/saira/OneDrive/Desktop/WordCount.py hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoop
WordCount.py                                100% 402      8.8KB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ scp -i C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem C:/Users/saira/OneDrive/Desktop/w.data hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoop
w.data                                       100% 528     13.3KB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$
```

Checking if the w.data file is in the cluster using `hadoop fs -ls/user/hadoop`

```
cp: '/home/hadoop/w.data': No such file or directory
[hadoop@ip-172-31-21-101 ~]$ hadoop fs -put /home/hadoop/w.data /user/hadoop
[hadoop@ip-172-31-21-101 ~]$ hadoop fs -ls /user/hadoop
Found 1 items
-rw-r--r--  1 hadoop hdfsadmingroup    528 2023-09-21 06:07 /user/hadoop/w.data
[hadoop@ip-172-31-21-101 ~]$
```

Successfully file is uploaded onto the cluster.

Now we run the python wordcount.py -r hadoop hdfs:///user/hadoop/w.data

This will take w.data as input data into the wordcount py and the MapReduce job will begin to run.

```
data
[hadoop@ip-172-31-21-101 ~]$ python WordCount.py -r hadoop hdfs:///user/hadoop/w
.data
No configs found; falling back on auto-configuration
No configs specified for hadoop runner
Looking for hadoop binary in $PATH...
Found hadoop binary: /usr/bin/hadoop
```

As we can see the output has words – word frequency

```
545746/output
Streaming final output from hdfs:///user/hadoop/tmp/mrjob/WordCount.hadoop.20230
921.060912.545746/output...
"a" 3
"all" 1
"an" 1
"and" 1
"are" 1
"as" 4
"available" 1
"be" 3
"by" 1
"cluster" 2
"combine" 1
"contained" 1
"defined" 1
"dependencies" 1
"do" 1
"either" 1
```

```
that" 1
the" 4
things" 1
those" 1
to" 3
two" 1
uploaded" 1
versions" 1
well" 1
when" 1
will" 1
within" 1
writing" 2
your" 5
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/WordCount.hadoop.202
921.060912.545746...
Removing temp directory /tmp/WordCount.hadoop.20230921.060912.545746...
[hadoop@ip-172-31-21-101 ~]$
```

Now we move the output into new directory

```
WRONG_REDUCE=0
job output is in hdfs:///user/hadoop/csp554
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/WordCount.hadoop.2023
0921.061246.096559...
Removing temp directory /tmp/WordCount.hadoop.20230921.061246.096559...
[hadoop@ip-172-31-21-101 ~]$
```

**Step 5.**

Word count2.py

Scp -i /path/to/pem/file /wordcount2.py

```
saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ scp -i 'C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem' 'C:/Users/saira/OneDrive/Desktop/WordCount2.py' hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoop
WordCount2.py                                100% 543    15.8kB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ |
```

6. WordCount2.py output here a\_to\_n 49 ,

Other 46

Output in new directory csp554

```
WRONG_REDUCE=0
job output is in hdfs:///user/hadoop/tmp/mrjob/WordCount2.hadoop.20230921.062222.992894/output
Streaming final output from hdfs:///user/hadoop/tmp/mrjob/WordCount2.hadoop.20230921.062222.992894/output...
"a_to_n"      49
"other"      46
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/WordCount2.hadoop.20230921.062222.992894...
Removing temp directory /tmp/WordCount2.hadoop.20230921.062222.992894...
[hadoop@ip-172-31-21-101 ~]$ |
```

Both programs are stored in /user/hadoop/csp554

```
[hadoop@ip-172-31-21-101 ~]$ hadoop fs -ls /user/hadoop/csp554
Found 2 items
-rw-r--r--  1 hadoop hdfsadmingroup      0 2023-09-21 06:13 /user/hadoop/csp554/_SUCCESS
-rw-r--r--  1 hadoop hdfsadmingroup    652 2023-09-21 06:13 /user/hadoop/csp554/part-00000
[hadoop@ip-172-31-21-101 ~]$ |
```

Salaries.py

```
saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ scp -i 'C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem' 'C:/Users/saira/OneDrive/Desktop/Salaries.py' hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoop
Salaries.py                                100% 411    11.1kB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ |
```



Upload salaries.tsv

```
saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ scp -i 'C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem' 'C:/Users/saira/OneDrive/Desktop/Salaries.tsv' hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoop
Salaries.tsv                                100% 1502KB   3.9MB/s   00:00
saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$
```

Step 7.

Moving tsv into home hadoop

```
cp: /home/hadoop/salaries.tsv: No such file or directory
[hadoop@ip-172-31-21-101 ~]$ hadoop fs -put /home/hadoop/Salaries.tsv /user/hadoop
[hadoop@ip-172-31-21-101 ~]$ hadoop fs -ls /user/hadoop
Found 4 items
-rw-r--r--  1 hadoop hdfsadmingroup 1538148 2023-09-21 06:34 /user/hadoop/Salaries.tsv
drwxr-xr-x  - hadoop hdfsadmingroup  0 2023-09-21 06:13 /user/hadoop/csp554
drwxr-xr-x  - hadoop hdfsadmingroup  0 2023-09-21 06:09 /user/hadoop/tmp
-rw-r--r--  1 hadoop hdfsadmingroup  528 2023-09-21 06:07 /user/hadoop/w.data
[hadoop@ip-172-31-21-101 ~]$
```

```
WRONG_REDUCE=0
job output is in hdfs:///user/hadoop/tmp/mrjob/Salaries.hadoop.20230921.064314.030329/output
Streaming final output from hdfs:///user/hadoop/tmp/mrjob/Salaries.hadoop.20230921.064314.030329/output...
"911 LEAD OPERATOR" 4
"911 OPERATOR SUPERVISOR" 4
"911 OPERATOR" 65
"ACCOUNT EXECUTIVE" 4
"ACCOUNTANT I" 15
"ACCOUNTANT II" 25
"ACCOUNTANT SUPV" 7
"ACCOUNTANT TRAINEE" 1
"ACCOUNTING ASST I" 6
"ACCOUNTING ASST II" 15
"ACCOUNTING ASST III" 33
"ACCOUNTING MANAGER" 2
"ACCOUNTING OPERATIONS OFFICER" 1
"ACCOUNTING SYSTEMS ADMINISTRAT" 3
"ACCOUNTING SYSTEMS ANALYST" 21
"ADM COORDINATOR" 2
"ADMINISTRATIVE AIDE, SHERIFF" 11
"ADMINISTRATIVE ANALYST I" 8
"ADMINISTRATIVE ANALYST II" 3
"ADMINISTRATIVE COORDINATOR" 10
"ADMINISTRATIVE POLICY ANALYST" 2
"ALCOHOL ASSESSMENT COUNSELOR I" 1
"ALCOHOL ASSESSMENT DIRECTOR CO" 1
"ALCOHOL ASSESSMT COUNSELOR II" 1
"ALCOHOL ASSESSMT COUNSELOR III" 1
"ANALYST/PROGRAMMER II" 6
```

9.) Salaries2.py

```
saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ scp -i 'C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem' 'C:/Users/saira/OneDrive/Desktop/Salaries2.py' hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoop
Salaries2.py                                100% 768    23.1KB/s   00:00
saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$
```

Salaries2.py with salaries.tsv

Python salaries2.py -r hadoop hdfs:///user/hadoop/salaries.tsv

execute the command above

```

ZONING EXAMINER I 2
"ZONING EXAMINER II" 1
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/Salaries.hadoop.20230921.064314.030329...
Removing temp directory /tmp/Salaries.hadoop.20230921.064314.030329...
[hadoop@ip-172-31-21-101 ~]$ python Salaries2.py -r hadoop hdfs:///user/hadoop/Salaries.tsv
No configs found; falling back on auto-configuration
No configs specified for hadoop runner
Looking for hadoop binary in $PATH...
Found hadoop binary: /usr/bin/hadoop
Using Hadoop version 2.10.1
Looking for Hadoop streaming jar in /home/hadoop/contrib...
Looking for Hadoop streaming jar in /usr/lib/hadoop-mapreduce...
Found Hadoop streaming jar: /usr/lib/hadoop-mapreduce/hadoop-streaming.jar
Creating temp directory /tmp/Salaries2.hadoop.20230921.064816.651188
uploading working dir files to hdfs:///user/hadoop/tmp/mrjob/Salaries2.hadoop.20230921.064816.651188/files/wd...

```

11). Count of employees output is on the terminal screen, that means our job is successful

```

CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
job output is in hdfs:///user/hadoop/tmp/mrjob/Salaries2.hadoop.20230921.064816.651188/output
Streaming final output from hdfs:///user/hadoop/tmp/mrjob/Salaries2.hadoop.20230921.064816.651188/output...
"High" 442
"Low" 7064
"Medium" 6312
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/Salaries2.hadoop.20230921.064816.651188...
Removing temp directory /tmp/Salaries2.hadoop.20230921.064816.651188...
[hadoop@ip-172-31-21-101 ~]$

```

12). Scp u.data

```

MINGW64/c/Users/saira/OneDrive/Desktop
e/Desktop/Salaries.py' hadoop@ec2-34-228-195-135.compute-1.amazonaws.com: /home/
hadoop
Salaries.py                                100% 411    11.1KB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
scp -i 'C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem' 'C:/Users/saira/OneDri
e/Desktop/Salaries.tsv' hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home
hadoop
Salaries.tsv                              100% 1502KB  3.9MB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
scp -i 'C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem' 'C:/Users/saira/OneDri
e/Desktop/Salaries2.py' hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home
hadoop
Salaries2.py                              100% 768    23.1KB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
scp -i 'C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem' 'C:/Users/saira/OneDri
e/Desktop/u.data' hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoo
.data                                     100% 2381KB  4.3MB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop

```



## 13).Scp movies.py

```
saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ scp -i 'C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem' 'C:/Users/saira/OneDrive/Desktop/u.data' hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoop
u.data                                100% 2381KB   4.3MB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ scp -i 'C:/Users/saira/OneDrive/Desktop/asn3-keypr.pem' 'C:/Users/saira/OneDrive/Desktop/Movies.py' hadoop@ec2-34-228-195-135.compute-1.amazonaws.com:/home/hadoop
Movies.py                             100% 492     14.8KB/s   00:00

saira@GhostBuster MINGW64 ~/OneDrive/Desktop
$ |
```

Hadoop fs -put /home/hadoop/u.data /user/hadoop

```
[hadoop@ip-172-31-21-101 ~]$ hadoop fs -put /home/hadoop/u.data /user/hadoop
[hadoop@ip-172-31-21-101 ~]$
```

Now execute python movies.py hadoop hdfs:///user/hadoop/u.data

```
[hadoop@ip-172-31-21-101 ~]$ python Movies5.py -r hadoop hdfs:///user/hadoop/u.data
[hadoop@ip-172-31-21-101 ~]$
```

python Movies.py -r hadoop hdfs:///user/hadoop/u.data

```
hadoop@ip-172-31-21-101:~
Connecting to Application History server at ip-172-31-21-101.ec2.internal/172.
1.21.101:10200
Connecting to ResourceManager at ip-172-31-21-101.ec2.internal/172.31.21.101:8
32
Connecting to Application History server at ip-172-31-21-101.ec2.internal/172.
1.21.101:10200
Loaded native gpl library
Successfully loaded & initialized native-lzo library [hadoop-lzo rev 049362b7c
53ff5f739d6b1532457f2c6cd495e8]
Total input files to process : 1
number of splits:4
Submitting tokens for job: job_1695275040009_0012
resource-types.xml not found
Unable to find 'resource-types.xml'.
Adding resource type - name = memory-mb, units = Mi, type = COUNTABLE
Adding resource type - name = vcores, units = , type = COUNTABLE
Submitted application application_1695275040009_0012
The url to track the job: http://ip-172-31-21-101.ec2.internal:20888/proxy/app
lication_1695275040009_0012/
Running job: job_1695275040009_0012
Job job_1695275040009_0012 running in uber mode : false
  map 0% reduce 0%
  map 50% reduce 0%
```

Output of the Movies job.

```
hadoop@ip-172-31-21-101:~
"81"    160
"82"    39
"83"    161
"84"    116
"85"    107
"86"    190
"87"    31
"88"    255
"89"    66
"9"     45
"90"    50
"91"    150
"92"    123
"93"    159
"94"    196
"95"    299
"96"    76
"97"    128
"98"    71
"99"    188
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/Moviesr.hadoop.20230921.152845.237872...
Removing temp directory /tmp/Moviesr.hadoop.20230921.152845.237872...
[hadoop@ip-172-31-21-101 ~]$
```

Finally Cluster terminated.

Amazon EMR > EMR on EC2: Clusters > asn3-cluster

asn3-cluster Updated less than a minute ago Actions

▼ Summary

Cluster info	Applications	Cluster management	Status and time
<p>Cluster ID</p> <p>j-1318LJ4U4MTB4</p> <p>Cluster configuration</p> <p>Instance groups</p> <p>Capacity</p> <p>1 Primary 1 Core 0 Task</p>	<p>Amazon EMR version</p> <p>emr-5.36.1</p> <p>Installed applications</p> <p>Hadoop 2.10.1, Hive 2.3.9, Hue 4.10.0, Mahout 0.13.0, Pig 0.17.0, Tez 0.9.2</p>	<p>Log destination in Amazon S3</p> <p><a href="#">aws-logs-965044767171-us-east-1/elasticmapreduce</a></p> <p>Persistent application UIs</p> <p><a href="#">YARN timeline server</a></p> <p><a href="#">Tez UI</a></p> <p>Primary node public DNS</p> <p><a href="#">ec2-34-228-195-135.compute-1.amazonaws.com</a></p> <p><a href="#">Connect to the Primary Node using SSH</a></p>	<p>Status</p> <p>Terminated</p> <p>Creation time</p> <p>September 21, 2023, 00:38 (UTC-05:00)</p> <p>Elapsed time</p> <p>9 hours, 52 minutes</p>

Properties Bootstrap actions Instances (Hardware) Steps Applications Configurations Monitoring Events Tags (0)

And resources are removed.

-----END-----